

Group-Reward Policy Optimization for Mathematical Reasoning

Minh Binh Le
University of Massachusetts Amherst
Amherst, MA
minble@umass.edu

October 2025

Abstract

Group-Reward Policy Optimization (GRPO) is a recent and efficient reinforcement learning algorithm for large language models alignment through relative reward comparison. In my project design, I explore the use of GRPO to enhance mathematical reasoning in decoder-only the base Qwen3-1.7B models. The training objective is to improve reasoning accuracy beyond 37% zero-shot baseline of the 1024 max tokens base model by applying GRPO and a GRPO \rightarrow SFT \rightarrow GRPO training method.

Keywords: GRPO · Reinforcement Fine-Tuning · Reasoning LLMs · Alignment

1 Introduction

Large language models, especially on-device language models, demonstrate strong performance on simple reasoning but are limited in long-chain reasoning. While supervised fine-tuning can improve some capabilities, the model does not generalize very well without sufficient high-quality data from SFT. Reinforcement learning methods like Proximal Policy Optimization (PPO) are inefficient as they require additional critic models to rank the rewards.

GRPO’s method makes it more efficient by comparing multiple rollouts within a batch to obtain better reasoning patterns. This is ideal for mathematical tasks where models can explore many different paths creatively that lead to the same result.

2 Background

2.1 Reinforcement Learning for LLM Alignment

Original RL methods on large language models like PPO need a critic model to reward responses based on the current fine-tuning models. This can be computationally expensive when you need to train an additional model to give the rewards and have two models running in parallel. GRPO, instead, uses token-level normalized advantages within rollout groups, enabling faster reward feedback without computational overheads.

2.2 Group-Reward Policy Optimization

$$\mathcal{L}_{\text{GRPO}}(\theta) = E_G \left[\frac{1}{n} \sum_{i=1}^n \min(r_i(\theta) A_i, \text{clip}(r_i(\theta), 1 - \epsilon, 1 + \epsilon) A_i) \right], \quad A_i = \frac{r_i - \mu_G}{\sigma_G + \delta}.$$

This formula explains that this policy is to get a smaller objective and make the model more stable while continuously reupdating the policy.

2.3 Experiments

I chose the dataset Countdown-Tasks-3to4 because it’s easy to experiment with and yields a diverse range of calculations the models need to compute and try out before answering, making it perfect for the experiment.

3 Implementation

3.1 Foundational Model

I used the open-source Qwen3-1.7B architecture in both 256 and 512 maximum token limits Both were trained using 8-bit quantization on a single NVIDIA A100 GPU.

3.2 Reward Function

Rewards were based on output correctness and format quality:

Condition	Reward
Correct numeric result	1.0
Near-correct	0.8–0.3
Well-formatted but wrong	0.1
Missing <answer> tag	0.0

Rewards were group-normalized before computing the policy loss. The per-token PPO loss used a clipping range of 0.25 to stabilize gradients.

3.3 Training

Two experimental setups were tested:

1. **Multi-stage GRPO → SFT → GRPO (256 max):** The first GRPO phase taught reasoning structure and tag compliance. The model self-improves by realizing it should cut down on random token generation and focus more on the math, allowing it to go from 0.2% to 37.3% accuracy in the first 80 RL steps.
SFT phase reinforced correct traces using curated examples. This allows the model to consolidate its previously explored correct reasoning trace and make training more stable
Final GRPO phase consolidated structured reasoning and made it more stable, allowing for further improvement by a further 10%, reaching the final accuracy of 48.8%
2. **Single-stage GRPO (512 max):** Longer context allowed extended reasoning but no supervised stabilization. It follows the same pattern as 256 max tokens generation. However, because it’s still exploring at the temperature of 1.0, the model generates some Laos and China reasoning tokens, making it less stable and achieving an improvement of up to 39.2% from the base of around 14%.

3.4 Computing Resources

Training ran for 160 steps for 256 max tokens with double GRPO and SFT, and only 80 steps for 512 max tokens took 40 minutes - 1.5 hours of compute on a single Nvidia A100.

4 Results

4.1 Performance Trends

The 256 model improved from 0% to 48.8% accuracy after the final GRPO phase. The 512 single-stage model reached 39.2%, showing that greater context alone did not guarantee stronger reasoning without SFT stabilization.

Model	Baseline (Zero-Shot)	Final Accuracy
GRPO \rightarrow SFT \rightarrow GRPO (256)	0.1%	48.8%
GRPO only (512)	15%	39.2%

4.2 Qualitative Behavior

The 256 model learned to produce structured expressions like $(79 - 60) + 17 = 36$ `<answer>36</answer>`. The 512 model performed longer chains but sometimes diverged by small margins, e.g., predicting 67 instead of 66 for multi-term operations.

5 Discussion

The experiments confirm that interleaving GRPO with SFT stabilizes learning. The SFT stage improves consistency, while GRPO reinforces exploration of reasoning variants.

My intuition for using GRPO was that relative comparison within a group can simulate “peer learning”: instead of a critic dictating correctness, the model self-organizes around relatively stronger trajectories. This approach rewards improvement rather than absolute success, leading to more resilient reasoning traces.

The larger 512 context helped retain more intermediate steps but also introduced noise—suggesting that reward shaping and intermediate supervision are crucial for stability.

6 Future Work

Future directions include:

- Applying double GRPO run symbolic reasoning.
- Experimenting with adaptive group sizes or hierarchical GRPO for longer reasoning horizons.
- Combining GRPO with retrieval-augmented or chain-of-thought supervision for compositional reasoning.

7 Conclusion

This study shows that Group-Reward Policy Optimization can enhance structured reasoning in compact language models. Even under limited compute, a multi-stage GRPO \rightarrow SFT \rightarrow GRPO framework enables small-context models to outperform larger zero-shot baselines by over 10 percentage points. By encouraging models to learn through relative self-comparison, GRPO provides a lightweight path toward autonomous reasoning improvement.

References

1. Schulman et al., *Proximal Policy Optimization Algorithms*, 2017.
2. OpenAI, *Group-Reward Policy Optimization for LLM Alignment*, 2024.
3. DeepSeek-R1, *Emergent Reasoning Behaviors via Group-Relative Reinforcement*, 2025.
4. Qwen3 Team, *Qwen3 1.7B Model Card and Training Details*, 2025.